

Vishal Sahil

RAG Knowledge Base

Canonical portfolio knowledge for AI chat and voice agents

Version 2.0 • Last updated 13 September 2026 • Built from current portfolio, project pages, repositories, and supplied resume/knowledge documents

0. RAG Usage Rules

- Treat this document as the canonical source for Vishal Sahil's public professional profile.
- Prefer current portfolio facts and current project pages when older resumes conflict with them.
- Use dated work-history entries for experience rather than exaggerating a single years-of-experience number.
- Keep answers concise by default and expand only when the visitor asks for technical depth.
- Never invent salary, age, date of birth, private address, confidential client details, or unlisted credentials.
- When a question is unrelated to Vishal or his professional work, explain that the assistant is designed to answer questions about Vishal.
- For project implementation details, source repositories may be newer than portfolio summaries; known differences are noted in this document.

1. Canonical Identity

Name: Vishal Sahil

Titles: AI Automation Engineer · AI Agents Builder · AI Developer · Prompt Engineer

Location: Karachi, Pakistan

Work Mode: Remote worldwide

Availability: Open to freelance, consulting, full-time, and relevant part-time opportunities

Email: vishalsahilofficial@gmail.com

Phone: +92-305-8377755

Linkedin: <https://linkedin.com/in/vishal-sahil-ai>

Portfolio: <https://vishalsahilai.vercel.app/>

Github: <https://github.com/vishalsahilai>

One-sentence introduction

Vishal Sahil is a Karachi-based AI Automation Engineer and AI Agents Builder who builds AI agents, RAG systems, voice/chat assistants, workflow automations, FastAPI backends, and full-stack AI products for businesses and recruiters evaluating his work.

Short professional summary

Vishal works at the intersection of LLMs, automation, and software engineering. His core toolkit includes Python, FastAPI, n8n, LangChain/LangGraph, React/Next.js, RAG with Pinecone, REST APIs, databases, and cloud deployment. He focuses on replacing repetitive operational work with reliable, measurable software and AI workflows.

2. What Vishal Does / Services

- AI agent and multi-agent system development
- RAG knowledge systems and document-grounded chatbots
- Real-time voice AI assistants
- n8n workflow automation and business process automation
- FastAPI backend/API development
- Full-stack AI application development
- Lead generation, enrichment, personalization, and outreach automation
- Third-party API integrations and workflow orchestration
- Prompt engineering, guardrails, fallback logic, and LLM evaluation
- Conversation memory and session architecture
- Automation reliability, retries, provider fallback, and API key rotation
- Responsive business websites and frontend development
- Google Sheets/Gmail/Slack-style operational integrations
- Testing, regression testing, and production deployment support

3. Complete Technical Stack

Languages & Frontend

Python, JavaScript, TypeScript, HTML5, CSS3, React, Next.js, Vite, Streamlit

Backend & APIs

FastAPI, Uvicorn, REST APIs, WebSockets, Pydantic, SQLAlchemy, Alembic, PyMySQL

AI / LLM / Agents

LangChain, LangGraph, AI Agents, Multi-Agent Systems, RAG, Conversational AI, Prompt Engineering, OpenAI API / GPT-4o, Google Gemini, Claude (Anthropic), Alibaba Cloud Qwen

RAG / Memory / Data

Pinecone, llama-text-embed-v2, HuggingFace embeddings, MongoDB / MongoDB Atlas, MySQL, Aiven, Supabase, Redis (documented production option in portfolio assistant)

Automation & Integrations

n8n, Make (Integromat), Zapier, Google Sheets API, Gmail / Gmail SMTP, Slack, Apify, Apollo, Workflow Orchestration, API Integration

Voice / Browser / Scraping

Deepgram Nova-3, ElevenLabs, Web Audio API, MediaSource, Playwright, BeautifulSoup, Requests, RapidFuzz, dateparser, python-dateutil, pypdf, python-docx

Deployment / Engineering

Vercel, Render, GitHub, CI/CD, Regression Testing, Loguru, Tenacity

4. Professional Experience

Artificial Intelligence Intern — FlyRank AI

Period: June 2026 – September 2026 · Part-time · Remote

- Developed 6 backend APIs using Python and FastAPI.
- Built authentication and 3 authorization systems using Supabase.
- Contributed to AI-powered backend workflows and intelligent applications.
- Worked with REST APIs, backend architecture, API validation, LLMs, RAG, and AI agents.
- Created and executed 16 automated tests for an AI backend evaluation suite.
- Participated in engineering discussions and solved practical backend-development tasks.

AI Automation Engineer — Freelance / Independent Consultant

Period: October 2025 – Present · Karachi, Pakistan / Remote

- Builds end-to-end n8n and AI automation workflows for founders and SMBs.
- Has integrated 10+ third-party APIs into automation systems.
- Documented result: reduced client onboarding time by more than 40% in a referenced workflow.
- Automates bottlenecks across sales, support, and operations.
- Designs reusable workflow templates and documented architectures so clients can operate and scale systems after delivery.
- Works directly with founders/product teams to scope, deploy, test, and iterate solutions.

Prompt Engineering / LLM Systems Experience

Earlier resumes also document Prompt Engineer work beginning in December 2025. The work includes advanced prompt frameworks, tool-integrated LLM agents, memory management, guardrails, provider/model fallback, structured-output reliability, regression testing, and evaluation across GPT, Claude, and Gemini-class models. In the current portfolio, prompt engineering is best treated as a core specialization used across projects rather than the only professional identity.

5. Selected Documented Outcomes

- 10+ hours of manual work saved per week for founders/SMBs (resume/knowledge-profile claim).
- 40%+ faster client onboarding in an automation engagement.
- 60% reduction in manual sales effort in the earlier lead-engagement resume description.
- 35%+ improvement in LLM task precision / structured-output reliability in the earlier prompt-engineering resume description.
- Automated content pipeline previously documented as scaling social output by 5×.
- Sadabahr resume version documents eliminating manual order logging and reducing order processing time from about five minutes to under 30 seconds per order.
- Customer-support resume version documents faster response handling and 24/7 automated coverage.
- These are historical/project-specific claims, not guarantees for every future engagement.

6. How Vishal Approaches Problems

1. Understand the real problem

Start with the workflow, users, bottlenecks, and desired outcome before choosing tools.

2. Build the smallest useful solution

Ship the narrowest version that solves the core problem correctly rather than overbuilding.

3. Test with real usage and data

Validate against real inputs, failures, and edge cases.

4. Improve and scale from results

Once the core flow works, improve reliability, automation depth, and scale based on evidence.

7. Projects — Canonical Project Knowledge

7.1 Portfolio AI Voice Agent — Sada

Year: 2026

Status: Live

Role: Sole engineer

Type: Real-time browser-based AI voice assistant

Aliases / retrieval terms: Sada, voice agent, portfolio voice assistant, AI voice assistant, real-time voice agent

What it is

A real-time voice assistant embedded in the portfolio. Visitors speak through the browser, see live transcription, and receive spoken answers grounded in Vishal's portfolio knowledge with conversation memory.

Problem it solves

A normal portfolio requires visitors to browse pages or type questions. Sada provides a natural voice interface while keeping answers grounded in real portfolio information.

How it works

Browser microphone audio is resampled to PCM16/16 kHz and streamed to FastAPI over WebSockets. Deepgram Nova-3 provides streaming speech-to-text. Portfolio context is retrieved from Pinecone, MongoDB stores multi-turn memory, the LLM streams an answer, and ElevenLabs streams speech back with synchronized text timing.

Key capabilities

- Live browser microphone streaming and interim transcription
- Deepgram Nova-3 streaming speech recognition
- Pinecone RAG with llama-text-embed-v2 embeddings and top semantic retrieval
- MongoDB multi-turn memory with a three-phase memory strategy and TTL cleanup
- Streaming LLM generation with provider fallback handling
- Current source repository documents Alibaba Cloud Qwen as primary LLM with Gemini fallback and key rotation; the portfolio project page still shows Gemini in the simplified flow
- ElevenLabs WebSocket streaming TTS with synchronized captions
- LISTENING → THINKING → SPEAKING call-state management
- Microphone suppression while the assistant is speaking
- Parallel retrieval, memory, and TTS preparation to reduce latency
- Browser-local daily voice usage limit documented in the repository: seven completed AI questions per browser per calendar day

Built with: Python, FastAPI, Uvicorn, WebSockets, Deepgram Nova-3, Pinecone, llama-text-embed-v2, MongoDB, Alibaba Cloud Qwen, Google Gemini fallback, ElevenLabs, Web Audio API, MediaSource, HTML, CSS, JavaScript

Portfolio page: <https://vishalsahilai.vercel.app/projects/portfolio-ai-voice-agent>

Repository: <https://github.com/vishalsahilai/portfolio-ai-voice-agent>

Live / demo: Embedded in the main portfolio website

7.2 Portfolio AI Assistant

Year: 2026

Status: Live / being upgraded to external RAG knowledge retrieval

Role: Sole engineer

Type: Portfolio conversational AI assistant

Aliases / retrieval terms: portfolio chatbot, portfolio assistant, AI assistant, portfolio bot, RAG chatbot

What it is

A conversational assistant embedded in the portfolio so recruiters, clients, and visitors can ask about skills, experience, projects, availability, and contact details without manually browsing the site.

Problem it solves

Visitors often want a direct answer about one skill, project, or contact detail. The assistant reduces navigation effort and provides conversational access to the portfolio.

How it works

A Next.js chat widget sends messages to a FastAPI /chat endpoint. Each visitor receives an isolated session ID. A three-phase hybrid memory system keeps prompts compact, LangChain orchestrates the LLM flow, and strict portfolio-only rules reduce hallucination. The current update is moving profile knowledge out of a static system prompt into an external RAG knowledge base such as this document.

Key capabilities

- FastAPI REST backend with POST /chat and health endpoint
- Session-isolated conversations to avoid cross-user leakage
- Three-phase hybrid memory with progressive summarization
- Maximum five stored summaries per session in the documented implementation
- Strict portfolio-only answering rules and input validation
- Graceful handling for empty, long, or failed LLM requests
- CORS integration with the Next.js portfolio frontend
- Structured logging and Render deployment
- Current migration target: RAG retrieval from indexed portfolio knowledge instead of keeping all profile facts in the system prompt

Built with: Python, FastAPI, LangChain, Google Gemini, RAG, Pydantic, Loguru, Next.js, React, REST API, Render; planned/updated knowledge layer uses vector retrieval

Portfolio page: <https://vishalsahilai.vercel.app/projects/portfolio-ai-assistant>

Repository: <https://github.com/vishalsahilai/portfolio-chatbot>

Live / demo: Embedded in the main portfolio website

7.3 Sadabahar Restaurant AI Chatbot

Year: 2025

Status: Live**Role:** Sole engineer**Type:** RAG-powered restaurant ordering AI agent**Aliases / retrieval terms:** Sadabahar chatbot, restaurant chatbot, restaurant ordering bot, menu RAG bot, restaurant AI agent

What it is

A production restaurant chatbot that answers menu questions from verified restaurant data, recognizes returning customers, takes orders conversationally, and automates order logging and confirmation.

Problem it solves

Restaurants repeatedly answer menu questions and manually process orders, while a generic LLM can hallucinate prices or menu items. The system needed grounded knowledge, customer recognition, reliable memory, and real operational actions after order confirmation.

How it works

FastAPI orchestrates the chat flow. Pinecone RAG with HuggingFace embeddings retrieves menu data before Gemini responds. MongoDB stores customers, sessions, and orders. A hybrid summarization memory controls context growth. Confirmed orders are saved, logged to Google Sheets, and followed by a branded HTML confirmation email.

Key capabilities

- Pinecone RAG over verified restaurant PDF/menu information
- HuggingFace all-MiniLM-L6-v2 embeddings
- Anti-hallucination safeguards for exact menu information
- Phone-based returning-customer recognition and reorder support
- Three-phase hybrid conversation memory and two-hour inactive-session expiry
- Customer/session isolation
- Automatic order storage in MongoDB and Google Sheets
- Branded HTML order confirmation emails through Gmail SMTP
- Gemini API rotation across multiple keys with fallback handling
- Smart RAG skip logic for simple order fields and confirmations

Built with: Python, FastAPI, LangChain, Google Gemini, HuggingFace, Pinecone, MongoDB Atlas, Google Sheets API, Gmail SMTP, Loguru, Tenacity, HTML, CSS, JavaScript**Portfolio page:** <https://vishalsahilai.vercel.app/projects/sadabahar-restaurant-ai-chatbot>**Repository:** <https://github.com/vishalsahilai/sadabahar-restaurant-chatbot>**Live / demo:** <https://sadabahar-restaurant-bot.vercel.app>

7.4 MediCore Hospital Management System

Year: 2026**Status:** Built / deployed**Role:** Full-stack developer**Type:** Full-stack hospital management platform**Aliases / retrieval terms:** MediCore, hospital management system, HMS, hospital dashboard, hospital app

What it is

A centralized hospital administration platform for managing doctors, patients, staff, admissions, shifts, and dashboard statistics through a responsive React interface backed by FastAPI and MySQL.

Problem it solves

Hospital records and staff operations become difficult to search and update when spread across disconnected processes. MediCore centralizes operational data and CRUD workflows in one application.

How it works

The React frontend communicates with a FastAPI REST API. Pydantic validates requests and responses, SQLAlchemy manages database interactions, and MySQL stores persistent records. Dashboard statistics are generated from live API data.

Key capabilities

- CRUD operations for doctors, patients, and staff
- Search and filtering on management pages
- Doctor specializations and doctor-to-patient assignment
- Patient admission tracking
- Morning, Evening, and Night staff shift management
- Live dashboard statistics and recent admissions
- SQLAlchemy ORM and Alembic database migration support
- Pydantic validation and Swagger/OpenAPI documentation
- Responsive frontend with cloud-hosted MySQL

Built with: Python, FastAPI, React, React Router, JavaScript, SQLAlchemy, Alembic, Pydantic, MySQL, PyMySQL, Aiven, Render, Vercel

Portfolio page: <https://vishalsahilai.vercel.app/projects/medicore-hms>

Repository: <https://github.com/vishalsahilai/hospital-management-system>

Live / demo: <https://medicore-hms-fastapi.vercel.app>

7.5 Multi-Agent Job Finder

Year: 2026

Status: Built

Role: Sole engineer

Type: AI-powered multi-board job discovery and ranking system

Aliases / retrieval terms: job finder, multi-agent job search, AI job finder, resume job matcher, job discovery agent

What it is

A job discovery system that reads an uploaded resume, generates optimized searches, finds real listings across multiple job boards, filters and deduplicates them, and ranks opportunities by resume relevance.

Problem it solves

Manual job searching requires repeated searches across job boards, date checking, duplicate removal, and manual comparison against a resume.

How it works

The application extracts PDF/DOCX resume text, identifies candidate skills and attributes with Python rules, uses Gemini once per run to analyze the profile and generate queries, collects job data with BeautifulSoup/Requests and Playwright, normalizes dates, removes duplicates, scores relevance, and presents ranked results in Streamlit.

Key capabilities

- PDF and DOCX resume parsing
- 100+ skill keyword dictionary and Python-based extraction
- Single Gemini API call per search run to reduce AI usage
- Search support for LinkedIn, Rozee.pk, Indeed, Glassdoor, Mustakbil, and Google Jobs
- Generation of multiple optimized job-board search URLs
- BeautifulSoup static scraping and Playwright dynamic-page support
- Junk URL filtering and multi-format date parsing
- Duplicate removal using normalized URLs and RapidFuzz matching
- Five-signal relevance scoring with percentage match ranking
- Ranked Streamlit job cards, application links, and CSV download
- Resume processed in memory without database storage in the documented design

Built with: Python, Streamlit, Google Gemini, pypdf, python-docx, BeautifulSoup, Requests, Playwright, dateparser, python-dateutil, RapidFuzz, Pydantic Settings, LangChain, Google GenAI

Portfolio page: <https://vishalsahilai.vercel.app/projects/multi-agent-job-finder>

Repository: <https://github.com/vishalsahilai/Multi-Agent-Job-Finder>

Live / demo: Portfolio project page links to the deployed/demo experience when available

7.6 AI Code Reviewer & Optimizer

Year: 2026

Status: Built

Role: AI / Full-Stack Engineer

Type: AI-powered Python code review and optimization platform

Aliases / retrieval terms: code reviewer, AI code optimizer, Python reviewer, LangGraph code reviewer, code audit agent

What it is

A developer tool that analyzes Python code, identifies security and quality problems, refactors the code, and generates technical documentation through a sequential multi-agent workflow.

Problem it solves

Manual code review can miss security, performance, maintainability, and style problems, and developers also spend time refactoring and documenting code after review.

How it works

A sequential LangGraph workflow sends Python code through specialized agents. The Scanner Agent produces a structured audit, the Refactor Agent improves the code based on the audit, and the Documentation Agent generates technical documentation.

Key capabilities

- Paste Python code or upload .py files

- Security vulnerability detection
- Code-quality and maintainability analysis
- Performance and style analysis
- Structured audit report and scoring
- AI-powered code refactoring
- Generated technical documentation
- Original-versus-refactored code comparison
- Pipeline progress/timing and responsive developer dashboard

Built with: Python, FastAPI, React, TypeScript, LangGraph, LangChain, Google Gemini, Vite

Portfolio page: <https://vishalsahilai.vercel.app/projects/ai-code-reviewer-optimizer>

Repository: <https://github.com/vishalsahilai/ai-code-reviewer>

Live / demo: Source code and demo media are available from the portfolio/project repository

7.7 AI-Powered Lead Engagement

Year: 2026

Status: Built

Role: Sole engineer

Type: Automated lead generation and personalized cold-email workflow

Aliases / retrieval terms: lead engagement system, lead generation automation, cold email automation, n8n lead workflow, outreach automation

What it is

An end-to-end n8n workflow that sources targeted businesses, removes duplicate leads, enriches missing email addresses, generates personalized outreach, stores campaign data, and sends cold emails automatically.

Problem it solves

Manual prospecting requires repetitive searching, contact copying, duplicate checks, email enrichment, message writing, and campaign tracking, which is slow and difficult to scale.

How it works

An n8n form collects campaign criteria. Apify sources business leads, existing records are checked for duplicates, Apollo enriches missing email addresses, qualified leads are processed with rate limiting, Gemini generates personalized outreach, Google Sheets stores the campaign data, and Gmail sends the messages in a loop until the campaign target is complete.

Key capabilities

- Campaign form for niche, country/location, and lead count
- Automated Apify lead sourcing and result retrieval
- Duplicate filtering and re-fetching when more unique leads are needed
- Apollo email enrichment for missing addresses
- One-by-one lead processing with rate limiting and delays
- Lead-specific Gemini personalization
- Google Sheets storage and separate cold-email tracking
- Automatic Gmail sending and campaign looping

Built with: n8n, Apify, Apollo, Google Gemini, Google Sheets, Gmail, REST APIs, JavaScript, JSON

Portfolio page: <https://vishalsahilai.vercel.app/projects/ai-powered-lead-engagement>

Repository: Not publicly listed on the portfolio project page

Live / demo: Workflow project; portfolio contains screenshots/demo information

7.8 Nexcore Client Website

Year: 2026

Status: Live

Role: Frontend Developer

Type: Responsive business website

Aliases / retrieval terms: Nexcore, client website, business website, vanilla JavaScript website, frontend client project

What it is

A clean, professional client website built from scratch with semantic HTML, custom CSS, and JavaScript, without React or another frontend framework.

Problem it solves

The client needed a simple professional online presence that clearly represented the business and worked well across desktop and mobile devices.

How it works

The site uses semantic HTML for structure, custom CSS for layout and responsive behavior, and JavaScript for browser interactions.

Key capabilities

- Custom website built from scratch
- Responsive desktop and mobile layout
- Semantic HTML structure
- Custom CSS and JavaScript interactions
- Lightweight implementation with no frontend framework dependency

Built with: HTML5, CSS3, JavaScript

Portfolio page: <https://vishalsahilai.vercel.app/projects/nexcore-client-website>

Repository: Not publicly listed on the portfolio project page

Live / demo: Live client website linked from the portfolio project page

8. Education & Certifications

Education

- Bachelor of Computer Science — Millennium Institute of Technology and Entrepreneurship (MITE), Karachi, Pakistan — 2024–2028 (expected).
- Computer Science (Intermediate) — Superior College of Science, Pakistan.
- Eva's Academy High School, Hyderabad — Computer Science — earlier profile lists September 2010 to September 2022.

Certifications / Training

- Agentic AI And Automation — Saylani Institute of Technology — 2026.
- Build Intelligent Agents Using DeepSeek & N8N — Coursera — skills highlighted: n8n and AI Agents.

9. Availability, Engagement, and Contact

- Current public availability: Remote · Freelance · Full-time.
- Earlier knowledge profile: Open to consulting and relevant part-time opportunities as well.
- Geography: Karachi, Pakistan; remote worldwide.
- Engagement model: Earlier knowledge document supports both fixed/project-based and hourly work depending on scope.
- Pricing: No fixed public price is listed. Discuss scope directly for a quote or compensation discussion.
- Typical delivery note: Earlier knowledge document says many automation/AI projects were delivered in roughly two-week cycles; actual timelines vary by scope.

Email: vishalsahilofficial@gmail.com

WhatsApp: +92-305-8377755

LinkedIn: <https://linkedin.com/in/vishal-sahil-ai>

Portfolio: <https://vishalsahilai.vercel.app/>

GitHub: <https://github.com/vishalsahilai>

The portfolio contact form sends the message directly from the website; the site states that it does not open an external email application.

10. References & Public Endorsements

- Nasir Hussain — Agentic AI Engineer @ Modulers Pvt Ltd. — teacher/mentor. Public reference emphasizes strong build discipline and attention to system quality.
- Haris Irfan — Senior SQA Automation Analyst @ Contour Software — mentor. Public reference emphasizes consistency and quality in Vishal's work.
- Muhammad Saad Naseem — Agentic AI Backend Engineer — teacher/mentor. Public reference describes Vishal as proactive, exceeding course expectations, building independent projects, and showing strong interest in innovation and learning.
- Awais — CEO, Nexcore — client. Public reference praises professionalism, an optimized and intuitive website, and on-schedule delivery.
- Dr. Vishal Dherwani — CEO, Ramsun Dental Care — client. Public reference expresses strong satisfaction with Google Business Profile optimization and digital-marketing work and recommends Vishal's services.

11. High-Priority Questions & Answers for RAG Retrieval

Each entry is deliberately self-contained so vector search can retrieve the answer even when a visitor asks a short or differently worded question.

FAQ-001 [HIGH] Who is Vishal Sahil?

Answer: Vishal Sahil is an AI Automation Engineer, AI Agents Builder, AI Developer, and prompt-engineering practitioner based in Karachi, Pakistan. He builds AI agents, workflow automations, RAG systems, voice/chat assistants, FastAPI backends, and full-stack AI products for businesses and recruiters evaluating his work.

Question variants / aliases: Who are you? | Tell me about Vishal | Introduce Vishal | What does Vishal do?

Keywords: bio, profile, AI automation engineer

FAQ-002 [HIGH] What does Vishal specialize in?

Answer: He specializes in AI automation, AI agents, RAG and conversational AI, n8n workflow automation, FastAPI backend development, API integrations, prompt engineering, multi-agent systems, and full-stack AI applications.

Question variants / aliases: What are his specialties? | Main expertise? | What is Vishal good at?

Keywords: specialization, expertise

FAQ-003 [HIGH] Is Vishal open to work?

Answer: Yes. His current portfolio states that he is open to work and available for remote freelance and full-time opportunities. His earlier knowledge profile also lists consulting and part-time opportunities when relevant.

Question variants / aliases: Are you available? | Can we hire Vishal? | Is Vishal looking for a job?

Keywords: availability, open to work, hiring

FAQ-004 [HIGH] Where is Vishal located?

Answer: Vishal is based in Karachi, Pakistan and works remotely with clients or teams worldwide.

Question variants / aliases: Where do you live? | What city is Vishal in? | Is Vishal in Pakistan?

Keywords: Karachi, Pakistan, remote

FAQ-005 [HIGH] How can I contact Vishal?

Answer: You can email him at vishalsahilofficial@gmail.com, connect on LinkedIn at [linkedin.com/in/vishal-sahil-ai](https://www.linkedin.com/in/vishal-sahil-ai), use the contact form on his portfolio, or message the public WhatsApp number +92-305-8377755.

Question variants / aliases: Contact Vishal | How do I reach you? | Email? | WhatsApp?

Keywords: contact, email, LinkedIn, WhatsApp

FAQ-006 [HIGH] What is Vishal's portfolio website?

Answer: His portfolio is <https://vishalsahil.ai>. It includes his profile, stack, experience, projects, references, education, contact form, text AI assistant, and voice AI assistant.

Question variants / aliases: Portfolio link? | Website?

Keywords: portfolio, website

FAQ-007 [HIGH] What is Vishal's GitHub?

Answer: His GitHub profile is <https://github.com/vishalsahilai>, where public repositories for several AI and full-stack projects are available.

Question variants / aliases: GitHub link? | Source code?

Keywords: GitHub, code, repositories

FAQ-008 [HIGH] What kind of work can Vishal build for a client?

Answer: He can build AI chatbots, voice agents, RAG systems, n8n automations, lead-generation and outreach workflows, FastAPI APIs, full-stack AI applications, multi-agent systems, API integrations, business websites, and custom internal tools.

Question variants / aliases: What services do you offer? | What can Vishal build? | Can he automate my business?

Keywords: services, client work, automation

FAQ-009 [HIGH] Does Vishal work with international clients?

Answer: Yes. His knowledge profile lists remote worldwide availability, and his portfolio is positioned for remote freelance and full-time work.

Question variants / aliases: Can he work with US clients? | Worldwide? | Remote?

Keywords: international, remote, worldwide

FAQ-010 [HIGH] How much experience does Vishal have?

Answer: His dated public experience includes freelance AI automation work from October 2025 onward and an AI internship at FlyRank AI from June to September 2026. An earlier knowledge profile described 2+ years when hands-on learning and preparation are included. For precise professional history, use the dated roles rather than overstating a single number.

Question variants / aliases: Years of experience? | How experienced is Vishal?

Keywords: experience, career

FAQ-011 [HIGH] What programming languages does Vishal use?

Answer: His main programming languages are Python and JavaScript, with TypeScript used in projects such as the AI Code Reviewer. He also works with HTML and CSS for frontend projects.

FAQ-012 [HIGH] Does Vishal know Python?

Answer: Yes. Python is his primary backend and AI-development language, used across FastAPI APIs, RAG systems, automation tools, job-search systems, and multi-agent applications.

FAQ-013 [HIGH] Does Vishal use FastAPI?

Answer: Yes. FastAPI is one of his core backend frameworks and is used in his portfolio assistant, Sada voice agent, Sadabahar restaurant chatbot, MediCore, and AI Code Reviewer.

FAQ-014 [HIGH] Does Vishal use n8n?

Answer: Yes. n8n is a core automation tool in his freelance work and projects, including lead sourcing, enrichment, personalized outreach, API orchestration, Google Sheets, and Gmail workflows.

FAQ-015 [HIGH] Does Vishal know LangChain?

Answer: Yes. He uses LangChain for LLM orchestration, conversational applications, RAG systems, and portfolio/chatbot projects.

FAQ-016 [HIGH] Does Vishal use LangGraph?

Answer: Yes. LangGraph is used in the AI Code Reviewer & Optimizer to implement a sequential multi-agent workflow.

FAQ-017 [HIGH] Does Vishal build RAG systems?

Answer: Yes. He has built RAG-based systems with Pinecone, including the Sadabahar restaurant chatbot and Sada voice agent, and is moving the portfolio AI assistant from static prompt knowledge to an external RAG knowledge base.

FAQ-018 [HIGH] Which vector database does Vishal use?

Answer: Pinecone is his main documented vector database. He has used it with HuggingFace embeddings in Sadabahar and Pinecone llama-text-embed-v2 in Sada.

FAQ-019 [HIGH] Which LLM providers has Vishal used?

Answer: His public work includes Google Gemini, OpenAI/GPT-4o, Claude, and Alibaba Cloud Qwen in the current Sada source architecture. Model choice depends on the project, cost, latency, and fallback requirements.

FAQ-020 [MEDIUM] Does Vishal build AI agents?

Answer: Yes. AI agents and multi-agent systems are a central part of his portfolio, including Sada, the code reviewer multi-agent pipeline, job discovery automation, and business workflow agents.

FAQ-021 [MEDIUM] Does Vishal build voice agents?

Answer: Yes. Sada is a real-time browser voice agent using streaming STT, RAG, conversation memory, streaming LLM responses, and TTS.

FAQ-022 [MEDIUM] Does Vishal build chatbots?

Answer: Yes. He has built portfolio and restaurant chatbots with session management, conversation memory, RAG/grounding, API integrations, and production deployments.

FAQ-023 [MEDIUM] Does Vishal know React and Next.js?

Answer: Yes. He uses React and Next.js for full-stack and portfolio interfaces, while also being able to build vanilla HTML/CSS/JavaScript websites such as Nexcore.

FAQ-024 [MEDIUM] Which databases has Vishal used?

Answer: His documented database experience includes MySQL, Aiven-hosted MySQL, MongoDB/MongoDB Atlas, Pinecone as a vector database, and Supabase for authentication/authorization work.

FAQ-025 [MEDIUM] Does Vishal work with REST APIs?

Answer: Yes. REST API design and third-party API integration are core parts of his work. He has built FastAPI endpoints and integrated services such as Gmail, Google Sheets, Apify, Apollo, Deepgram, ElevenLabs, Gemini, and others.

FAQ-026 [MEDIUM] Does Vishal have testing experience?

Answer: Yes. At FlyRank AI he created and executed 16 automated tests for an AI backend evaluation suite, and his resumes also highlight regression testing and model/prompt evaluation work.

FAQ-027 [MEDIUM] Does Vishal know CI/CD and deployment?

Answer: Yes. His stack includes CI/CD concepts and deployments across Vercel, Render, Aiven, and GitHub-based workflows.

FAQ-028 [MEDIUM] Can Vishal integrate Gmail and Google Sheets?

Answer: Yes. He has used Gmail and Google Sheets in production-style automations for restaurant order confirmation and lead-engagement workflows.

FAQ-029 [MEDIUM] Does Vishal use Make or Zapier?

Answer: Yes. Make (Integromat) and Zapier are part of his automation toolkit, alongside n8n.

FAQ-030 [MEDIUM] Does Vishal use Playwright or web scraping tools?

Answer: Yes. The Multi-Agent Job Finder uses Playwright for dynamic pages and BeautifulSoup/Requests for static scraping.

FAQ-031 [HIGH] What did Vishal do at FlyRank AI?

Answer: From June to September 2026, Vishal worked part-time and remotely as an Artificial Intelligence Intern at FlyRank AI. He developed six backend APIs with Python/FastAPI, built authentication and three authorization systems using Supabase, contributed to AI-powered backend workflows, created and ran 16 automated tests, and applied LLMs, RAG, and AI agents in backend applications.

FAQ-032 [HIGH] What is Vishal's freelance experience?

Answer: Since October 2025, he has worked as a freelance AI Automation Engineer/independent consultant, building end-to-end n8n and AI workflows, integrating 10+ third-party APIs, improving onboarding and repetitive business processes, and creating reusable automation architectures.

FAQ-033 [HIGH] Has Vishal worked as a Prompt Engineer?

Answer: Prompt engineering is a documented part of his freelance background. Earlier resumes list a Prompt Engineer role from December 2025 onward, covering tool-integrated LLM agents, memory, guardrails, fallback handling, prompt regression testing, and evaluation across multiple models.

FAQ-034 [HIGH] What measurable results are documented in Vishal's resumes?

Answer: His supplied resumes document outcomes including 40% faster client onboarding, 60% lower manual sales effort in a lead workflow, 35%+ improvement in structured LLM task precision, and more than 10 hours of manual work saved per week for founders/SMBs. Treat these as project/resume claims rather than universal guarantees for every engagement.

FAQ-035 [HIGH] How does Vishal approach a new problem?

Answer: His public process is: understand the real workflow and bottlenecks, build the smallest useful solution, test it with real usage and edge cases, then improve reliability and scale based on results.

FAQ-036 [HIGH] What is Portfolio AI Voice Agent — Sada?

Answer: A real-time voice assistant embedded in the portfolio. Visitors speak through the browser, see live transcription, and receive spoken answers grounded in Vishal's portfolio knowledge with conversation memory.

Question variants / aliases: Sada | voice agent | portfolio voice assistant | AI voice assistant | real-time voice agent

Keywords: Portfolio AI Voice Agent — Sada

FAQ-037 [MEDIUM] How does Portfolio AI Voice Agent — Sada work?

Answer: Browser microphone audio is resampled to PCM16/16 kHz and streamed to FastAPI over WebSockets. Deepgram Nova-3 provides streaming speech-to-text. Portfolio context is retrieved from Pinecone, MongoDB stores multi-turn memory, the LLM streams an answer, and ElevenLabs streams speech back with synchronized text timing.

Question variants / aliases: Sada | voice agent | portfolio voice assistant | AI voice assistant | real-time voice agent

Keywords: Portfolio AI Voice Agent — Sada, architecture, workflow

FAQ-038 [MEDIUM] What technologies are used in Portfolio AI Voice Agent — Sada?

Answer: Python, FastAPI, Uvicorn, WebSockets, Deepgram Nova-3, Pinecone, llama-text-embed-v2, MongoDB, Alibaba Cloud Qwen, Google Gemini fallback, ElevenLabs, Web Audio API, MediaSource, HTML, CSS, JavaScript

Question variants / aliases: Sada | voice agent | portfolio voice assistant | AI voice assistant | real-time voice agent

Keywords: Portfolio AI Voice Agent — Sada, tech stack

FAQ-039 [HIGH] What is Portfolio AI Assistant?

Answer: A conversational assistant embedded in the portfolio so recruiters, clients, and visitors can ask about skills, experience, projects, availability, and contact details without manually browsing the site.

Question variants / aliases: portfolio chatbot | portfolio assistant | AI assistant | portfolio bot | RAG chatbot

Keywords: Portfolio AI Assistant

FAQ-040 [MEDIUM] How does Portfolio AI Assistant work?

Answer: A Next.js chat widget sends messages to a FastAPI /chat endpoint. Each visitor receives an isolated session ID. A three-phase hybrid memory system keeps prompts compact, LangChain orchestrates the LLM flow, and strict portfolio-only rules reduce hallucination. The current update is moving profile knowledge out of a static system prompt into an external RAG knowledge base such as this document.

Question variants / aliases: portfolio chatbot | portfolio assistant | AI assistant | portfolio bot | RAG chatbot

Keywords: Portfolio AI Assistant, architecture, workflow

FAQ-041 [MEDIUM] What technologies are used in Portfolio AI Assistant?

Answer: Python, FastAPI, LangChain, Google Gemini, RAG, Pydantic, Loguru, Next.js, React, REST API, Render; planned/updated knowledge layer uses vector retrieval

Question variants / aliases: portfolio chatbot | portfolio assistant | AI assistant | portfolio bot | RAG chatbot

Keywords: Portfolio AI Assistant, tech stack

FAQ-042 [HIGH] What is Sadabahar Restaurant AI Chatbot?

Answer: A production restaurant chatbot that answers menu questions from verified restaurant data, recognizes returning customers, takes orders conversationally, and automates order logging and confirmation.

Question variants / aliases: Sadabahar chatbot | restaurant chatbot | restaurant ordering bot | menu RAG bot | restaurant AI agent

Keywords: Sadabahar Restaurant AI Chatbot

FAQ-043 [MEDIUM] How does Sadabahar Restaurant AI Chatbot work?

Answer: FastAPI orchestrates the chat flow. Pinecone RAG with HuggingFace embeddings retrieves menu data before Gemini responds. MongoDB stores customers, sessions, and orders. A hybrid summarization memory controls context growth. Confirmed orders are saved, logged to Google Sheets, and followed by a branded HTML confirmation email.

Question variants / aliases: Sadabahar chatbot | restaurant chatbot | restaurant ordering bot | menu RAG bot | restaurant AI agent

Keywords: Sadabahar Restaurant AI Chatbot, architecture, workflow

FAQ-044 [MEDIUM] What technologies are used in Sadabahar Restaurant AI Chatbot?

Answer: Python, FastAPI, LangChain, Google Gemini, HuggingFace, Pinecone, MongoDB Atlas, Google Sheets API, Gmail SMTP, Loguru, Tenacity, HTML, CSS, JavaScript

Question variants / aliases: Sadabahar chatbot | restaurant chatbot | restaurant ordering bot | menu RAG bot | restaurant AI agent

Keywords: Sadabahar Restaurant AI Chatbot, tech stack

FAQ-045 [HIGH] What is MediCore Hospital Management System?

Answer: A centralized hospital administration platform for managing doctors, patients, staff, admissions, shifts, and dashboard statistics through a responsive React interface backed by FastAPI and MySQL.

Question variants / aliases: MediCore | hospital management system | HMS | hospital dashboard | hospital app

Keywords: MediCore Hospital Management System

FAQ-046 [MEDIUM] How does MediCore Hospital Management System work?

Answer: The React frontend communicates with a FastAPI REST API. Pydantic validates requests and responses, SQLAlchemy manages database interactions, and MySQL stores persistent records. Dashboard statistics are generated from live API data.

Question variants / aliases: MediCore | hospital management system | HMS | hospital dashboard | hospital app

Keywords: MediCore Hospital Management System, architecture, workflow

FAQ-047 [MEDIUM] What technologies are used in MediCore Hospital Management System?

Answer: Python, FastAPI, React, React Router, JavaScript, SQLAlchemy, Alembic, Pydantic, MySQL, PyMySQL, Aiven, Render, Vercel

Question variants / aliases: MediCore | hospital management system | HMS | hospital dashboard | hospital app

Keywords: MediCore Hospital Management System, tech stack

FAQ-048 [HIGH] What is Multi-Agent Job Finder?

Answer: A job discovery system that reads an uploaded resume, generates optimized searches, finds real listings across multiple job boards, filters and deduplicates them, and ranks opportunities by resume relevance.

Question variants / aliases: job finder | multi-agent job search | AI job finder | resume job matcher | job discovery agent

Keywords: Multi-Agent Job Finder

FAQ-049 [MEDIUM] How does Multi-Agent Job Finder work?

Answer: The application extracts PDF/DOCX resume text, identifies candidate skills and attributes with Python rules, uses Gemini once per run to analyze the profile and generate queries, collects job data with BeautifulSoup/Requests and Playwright, normalizes dates, removes duplicates, scores relevance, and presents ranked results in Streamlit.

Question variants / aliases: job finder | multi-agent job search | AI job finder | resume job matcher | job discovery agent

Keywords: Multi-Agent Job Finder, architecture, workflow

FAQ-050 [MEDIUM] What technologies are used in Multi-Agent Job Finder?

Answer: Python, Streamlit, Google Gemini, pypdf, python-docx, BeautifulSoup, Requests, Playwright, dateparser, python-dateutil, RapidFuzz, Pydantic Settings, LangChain, Google GenAI

Question variants / aliases: job finder | multi-agent job search | AI job finder | resume job matcher | job discovery agent

Keywords: Multi-Agent Job Finder, tech stack

FAQ-051 [HIGH] What is AI Code Reviewer & Optimizer?

Answer: A developer tool that analyzes Python code, identifies security and quality problems, refactors the code, and generates technical documentation through a sequential multi-agent workflow.

Question variants / aliases: code reviewer | AI code optimizer | Python reviewer | LangGraph code reviewer | code audit agent

Keywords: AI Code Reviewer & Optimizer

FAQ-052 [MEDIUM] How does AI Code Reviewer & Optimizer work?

Answer: A sequential LangGraph workflow sends Python code through specialized agents. The Scanner Agent produces a structured audit, the Refactor Agent improves the code based on the audit, and the Documentation Agent generates technical documentation.

Question variants / aliases: code reviewer | AI code optimizer | Python reviewer | LangGraph code reviewer | code audit agent

Keywords: AI Code Reviewer & Optimizer, architecture, workflow

FAQ-053 [MEDIUM] What technologies are used in AI Code Reviewer & Optimizer?

Answer: Python, FastAPI, React, TypeScript, LangGraph, LangChain, Google Gemini, Vite

Question variants / aliases: code reviewer | AI code optimizer | Python reviewer | LangGraph code reviewer | code audit agent

Keywords: AI Code Reviewer & Optimizer, tech stack

FAQ-054 [HIGH] What is AI-Powered Lead Engagement?

Answer: An end-to-end n8n workflow that sources targeted businesses, removes duplicate leads, enriches missing email addresses, generates personalized outreach, stores campaign data, and sends cold emails automatically.

Question variants / aliases: lead engagement system | lead generation automation | cold email automation | n8n lead workflow | outreach automation

Keywords: AI-Powered Lead Engagement

FAQ-055 [MEDIUM] How does AI-Powered Lead Engagement work?

Answer: An n8n form collects campaign criteria. Apify sources business leads, existing records are checked for duplicates, Apollo enriches missing email addresses, qualified leads are processed with rate limiting, Gemini generates personalized outreach, Google Sheets stores the campaign data, and Gmail sends the messages in a loop until the campaign target is complete.

Question variants / aliases: lead engagement system | lead generation automation | cold email automation | n8n lead workflow | outreach automation

Keywords: AI-Powered Lead Engagement, architecture, workflow

FAQ-056 [MEDIUM] What technologies are used in AI-Powered Lead Engagement?

Answer: n8n, Apify, Apollo, Google Gemini, Google Sheets, Gmail, REST APIs, JavaScript, JSON

Question variants / aliases: lead engagement system | lead generation automation | cold email automation | n8n lead workflow | outreach automation

Keywords: AI-Powered Lead Engagement, tech stack

FAQ-057 [HIGH] What is Nexcore Client Website?

Answer: A clean, professional client website built from scratch with semantic HTML, custom CSS, and JavaScript, without React or another frontend framework.

Question variants / aliases: Nexcore | client website | business website | vanilla JavaScript website | frontend client project

Keywords: Nexcore Client Website

FAQ-058 [MEDIUM] How does Nexcore Client Website work?

Answer: The site uses semantic HTML for structure, custom CSS for layout and responsive behavior, and JavaScript for browser interactions.

Question variants / aliases: Nexcore | client website | business website | vanilla JavaScript website | frontend client project

Keywords: Nexcore Client Website, architecture, workflow

FAQ-059 [MEDIUM] What technologies are used in Nexcore Client Website?

Answer: HTML5, CSS3, JavaScript

Question variants / aliases: Nexcore | client website | business website | vanilla JavaScript website | frontend client project

Keywords: Nexcore Client Website, tech stack

FAQ-060 [MEDIUM] What is Vishal studying?

Answer: He is pursuing a Bachelor of Computer Science at the Millennium Institute of Technology and Entrepreneurship (MITE), Karachi, with an expected 2024–2028 study period.

FAQ-061 [MEDIUM] Where did Vishal study before university?

Answer: His earlier profile lists Computer Science (Intermediate) at Superior College of Science and Computer Science schooling at Eva's Academy High School in Hyderabad.

FAQ-062 [MEDIUM] What certifications does Vishal have?

Answer: His portfolio lists Agentic AI And Automation from Saylani Institute of Technology (2026) and a Coursera credential titled "Build Intelligent Agents Using DeepSeek & N8N," focused on n8n and AI agents.

FAQ-063 [MEDIUM] What are Vishal's pricing or rates?

Answer: No fixed public rate is listed. His earlier knowledge document says he can work project-based or hourly depending on scope, and specific pricing should be discussed directly based on the project.

FAQ-064 [MEDIUM] How long does a typical project take?

Answer: An earlier knowledge document states that many automation/AI projects were delivered in roughly two-week cycles. Actual timing depends on scope, integrations, testing, and deployment requirements.

FAQ-065 [MEDIUM] Can Vishal work full-time?

Answer: Yes. The current portfolio explicitly lists availability for remote freelance and full-time work.

FAQ-066 [MEDIUM] Can Vishal work part-time or as a consultant?

Answer: His earlier knowledge profile lists consulting and part-time work as possible engagement types, depending on the opportunity and scope.

FAQ-067 [MEDIUM] Can I message Vishal through the portfolio?

Answer: Yes. The portfolio includes a contact form that sends the message directly; the site states that it does not open a separate email application.

FAQ-068 [MEDIUM] Who has recommended Vishal?

Answer: The portfolio includes references from mentors and clients: Nasir Hussain, Haris Irfan, Muhammad Saad Naseem, Awais (CEO of Nexcore), and Dr. Vishal Dherwani (CEO of Ramsun Dental Care). Their comments emphasize proactivity, delivery quality, professionalism, consistency, and client satisfaction.

FAQ-069 [MEDIUM] Has Vishal done client website work?

Answer: Yes. Nexcore is a live client website built from scratch with HTML, CSS, and JavaScript. A public client reference from Nexcore's CEO praises the site's professionalism, optimization, usability, and on-schedule completion.

FAQ-070 [MEDIUM] Has Vishal done digital marketing or Google Business Profile work?

Answer: Yes. His earlier skills include digital marketing, and the portfolio includes a client reference from Ramsun Dental Care praising his Google Business Profile optimization and digital marketing work.

FAQ-071 [MEDIUM] Can Vishal build an AI chatbot for my business?

Answer: Yes. His portfolio includes both a portfolio AI assistant and a restaurant chatbot with memory, grounding/RAG, integrations, and deployment. The exact architecture would depend on your data and workflow.

FAQ-072 [MEDIUM] Can Vishal build a RAG chatbot from my PDFs or documents?

Answer: Yes. Sadabahar retrieves verified restaurant information from documents with Pinecone RAG, and his current portfolio assistant update is specifically moving portfolio knowledge into a RAG document store.

FAQ-073 [MEDIUM] Can Vishal build a voice AI assistant?

Answer: Yes. Sada demonstrates real-time browser audio streaming, Deepgram STT, RAG, memory, LLM streaming, and ElevenLabs TTS. Similar patterns can be adapted to other voice-assistant use cases.

FAQ-074 [MEDIUM] Can Vishal automate lead generation and cold outreach?

Answer: Yes. His AI-Powered Lead Engagement workflow automates business lead sourcing, deduplication, email enrichment, AI personalization, tracking, and Gmail outreach through n8n.

FAQ-075 [MEDIUM] Can Vishal build backend APIs?

Answer: Yes. He has built FastAPI APIs across multiple projects and developed six backend APIs during his FlyRank AI internship.

FAQ-076 [MEDIUM] Can Vishal build a full-stack application?

Answer: Yes. MediCore is a full-stack React/FastAPI/MySQL application, and several other projects combine backend AI services with web frontends.

FAQ-077 [MEDIUM] Can Vishal build an internal automation tool?

Answer: Yes. His freelance work focuses on eliminating repetitive operations with AI workflows, API integrations, n8n, and custom software.

FAQ-078 [MEDIUM] Can Vishal build a multi-agent system?

Answer: Yes. The AI Code Reviewer uses a sequential multi-agent workflow, and multi-agent systems are listed among his core capabilities.

FAQ-079 [MEDIUM] Can Vishal optimize an existing automation or AI agent?

Answer: Yes. His work includes workflow bottleneck analysis, prompt improvement, fallback handling, testing, API rotation, and reliability improvements.

FAQ-080 [MEDIUM] Can Vishal create a business website without React?

Answer: Yes. Nexcore was built from scratch using HTML, CSS, and JavaScript without React or Next.js.

FAQ-081 [GUARDRAIL] What is Vishal's age or date of birth?

Answer: This information is not included in the public portfolio or the supplied professional documents. Do not guess.

FAQ-082 [GUARDRAIL] What salary does Vishal expect?

Answer: No salary expectation is publicly specified. For a job opportunity, ask Vishal directly and provide the role, location, employment type, and compensation range.

FAQ-083 [GUARDRAIL] What is Vishal's exact hourly rate?

Answer: No exact public hourly rate is specified. Pricing depends on scope and should be discussed directly.

FAQ-084 [GUARDRAIL] What is Vishal's home address?

Answer: Only Karachi, Pakistan is provided publicly. Do not provide or infer a street address.

FAQ-085 [GUARDRAIL] Is Vishal married or what is his personal family status?

Answer: This is not part of the professional portfolio or knowledge base. Do not guess or infer it.

FAQ-086 [GUARDRAIL] What private clients has Vishal worked with?

Answer: Only publicly named client/project information in this knowledge base should be shared. Do not invent or reveal confidential client information.

FAQ-087 [GUARDRAIL] Does Vishal guarantee a specific percentage improvement?

Answer: No. Historical project metrics describe past outcomes or resume claims, not guarantees for future projects. Results depend on the workflow, data quality, integrations, and scope.

FAQ-088 [GUARDRAIL] Can the AI assistant answer questions unrelated to Vishal?

Answer: The portfolio assistant should stay focused on Vishal's skills, experience, projects, services, availability, and contact information. For unrelated questions, it should politely say it is designed to answer questions about Vishal.

12. Retrieval Alias Index

Identity: Vishal Sahil, Vishal, Sahil, AI Automation Engineer, AI Developer, AI Agents Builder, Prompt Engineer

Voice: Sada, Portfolio AI Voice Agent, voice agent, voice assistant, Deepgram, ElevenLabs

Portfolio Chatbot: Portfolio AI Assistant, portfolio chatbot, portfolio bot, RAG assistant

Restaurant: Sadabahar, restaurant chatbot, restaurant ordering AI, menu RAG

Hospital: MediCore, HMS, hospital management system, hospital dashboard

Jobs: Multi-Agent Job Finder, job finder, resume job matcher, AI job search

Code Review: AI Code Reviewer, code optimizer, Python audit, LangGraph reviewer

Lead Automation: AI-Powered Lead Engagement, lead generation, cold email automation, Apify, Apollo

Website: Nexcore, client website, business website, vanilla JavaScript website

Automation: n8n, workflow automation, business automation, API integration

RAG: Pinecone, retrieval augmented generation, vector search, embeddings, knowledge base

13. Do-Not-Hallucinate / Unknown Information

- Age and date of birth are not provided publicly.
- Exact salary expectation is not provided publicly.
- Exact hourly rate is not provided publicly.
- Street/home address is not public; only Karachi, Pakistan should be used.
- Personal family or relationship status is not part of this professional knowledge base.
- Do not disclose or invent confidential clients, private project data, API keys, passwords, or internal credentials.
- Do not claim every project achieved the same historical percentage improvements.
- If a project detail is not in this document, say it is not specified rather than guessing.

14. Recommended Agent Response Behavior

- Answer directly in 1–4 sentences for ordinary visitor questions.
- For technical project questions, provide the project purpose, architecture/flow, and stack, then offer more detail if requested.
- For hiring questions, mention availability and give contact methods.
- For “best project” or subjective comparisons, explain that the best example depends on the use case: Sada for real-time voice AI, Sadabahar for RAG + operations, MediCore for full-stack CRUD, Lead Engagement for n8n automation, Job Finder for data/search pipelines, and Code Reviewer for multi-agent developer tooling.
- When the user asks for source code, provide the relevant public repository if listed.
- When the user asks for a live project, provide the listed live link or portfolio project page.
- Never fabricate a missing credential, client, result, price, or personal detail.

15. Source Links / Maintenance References

- <https://vishalsahilai.vercel.app/>
- <https://github.com/vishalsahilai>

- <https://github.com/vishalsahilai/portfolio-ai-voice-agent>
- <https://github.com/vishalsahilai/portfolio-chatbot>
- <https://github.com/vishalsahilai/sadabahar-restaurant-chatbot>
- <https://github.com/vishalsahilai/hospital-management-system>
- <https://github.com/vishalsahilai/Multi-Agent-Job-Finder>
- <https://github.com/vishalsahilai/ai-code-reviewer>

Maintenance note: when the portfolio changes, update the canonical identity/experience/project sections first, then update or add FAQ entries so the RAG index contains both the facts and the likely natural-language questions visitors will ask.